

Build an AI Analyst You Can Trust

BrainStation Toronto · Khaled Zaky · the one rule, the nine controls, the four questions

Every fact has to carry the exact sentence it came from.
Something has to go back and check that the sentence is really there.

THE NINE CONTROLS

WHAT GOES IN	WHAT COMES OUT	WHAT PROVES IT
01 Source of record every input traces back to where it came from	04 Structured artifact Problem, Evidence, Options, Recommendation, Ask	07 Evidence pack the receipt: each claim, its source, and whether it passed
02 Grounding contract every claim carries the exact sentence it came from	05 Groundedness check ordinary code re-checks each claim against its source	08 Run metrics groundedness, completion, cost per run, latency
03 Scope boundary the agent knows what it may and may not touch	06 Fail-closed gate a failed check stops the brief from shipping	09 Regression set the same checks rerun on every change

Engineers call this an eval, risk teams call it a guardrail, and regulators call it governance. It's the same code.

THE FOUR QUESTIONS

- 1 Can every fact name its source?
- 2 Does something check that source automatically, or only a human on a good day?
- 3 When the check fails, does it block the document, or warn and get ignored?
- 4 Could you show someone the receipt six months from now?

If any answer is no, you've built a demo that isn't ready to be a teammate yet.

THE FOUR NUMBERS (TONIGHT'S BRIEF)

- Groundedness:** 3 of 4 on the first run, 3 of 3 after the fix
- Completion:** 5 of 5 brief sections present
- Cost per run:** about \$0.08 a run (estimated)
- Latency:** about 19 seconds a run (estimated)



Rebuild it with no code, or with Claude Code
khaledzaky.com/trust
repo: github.com/kzaky/freshcart-agent



Go deeper on evals
khaledzaky.com/trust/evals

FreshCart, Provision Retail Analytics, InstaFresh and Grocerly are fictional. The failure mode is real.