

Build an AI Analyst You Can Trust

BrainStation Toronto · Khaled Zaky · the one rule, the nine controls, the four questions

Every fact has to carry the exact sentence it came from.
Something has to go back and check that the sentence is really there.

THE NINE CONTROLS

WHAT GOES IN	WHAT COMES OUT	WHAT PROVES IT
01 Source of record every input traceable to where it came from	04 Structured artifact Problem, Evidence, Options, Recommendation, Ask	07 Evidence pack the receipt: every claim, its source, its result
02 Grounding contract every claim carries the exact sentence it came from	05 Groundedness check every claim re-verified against its source by code	08 Run metrics groundedness, completion, cost per run, latency
03 Scope boundary the agent knows what it may and may not touch	06 Fail-closed gate a failed check blocks — it does not warn	09 Regression set the same checks run every time, not once

Your engineer calls 04–06 an eval. Your risk team calls them a guardrail. Your regulator calls 07–09 governance. Same code.

THE FOUR QUESTIONS

- 1 Can every fact name its source?
- 2 Does something automatically check that source, or just a human on a good day?
- 3 When the check fails, does it block, or does it warn and get ignored?
- 4 Can you show someone the receipt six months from now?

If any answer is no, you have a demo. Not yet a teammate.

THE FOUR NUMBERS (TONIGHT)

- Groundedness.** 3 of 4, one caught.
Completion. all five brief sections present
Cost per run. about \$0.11
Latency. 19 seconds



Rebuild it — no code, or with the repo
khaledzaky.com/trust



Go deeper on evals
khaledzaky.com/trust/evals

FreshCart, Provision Retail Analytics, InstaFresh and Grocerly are fictional. The failure mode is not.