

Build an AI Analyst You Can Trust

BrainStation Toronto · Khaled Zaky · the one rule, the nine controls, the four questions

Every fact has to carry the exact sentence it came from.

Something has to go back and check that the sentence is really there.

The nine controls

What goes in

01 Source of record

every input is traceable to where it came from

02 Grounding contract

every claim carries the exact sentence it came from

03 Scope boundary

the agent knows what it may and may not touch

What comes out

04 Structured artifact

Problem, Evidence, Options, Recommendation, Ask

05 Groundedness check

every claim re-verified against its source by ordinary code

06 Fail-closed gate

a failed check blocks. it does not warn

What proves it

07 Evidence pack

the receipt: every claim, its source, its result

08 Run metrics

groundedness, completion, cost per run, latency

09 Regression set

the same checks run every time, not once

Your engineer calls 04 to 06 an eval. Your risk team calls them a guardrail. Your regulator calls 07 to 09 governance. Same code.

The four questions, before any AI output reaches a real decision

1. Can every fact name its source?
2. Does something automatically check that source, or just a human on a good day?
3. When the check fails, does it block, or does it warn and get ignored?
4. Can you show someone the receipt six months from now?

If any answer is no, you have a demo. Not yet a teammate.

The four numbers that are actually a number

Groundedness: how many cited claims held up against their source. Tonight: 3 of 4, one caught.

Completion: did it produce the thing in the required shape.

Cost per run: what one decision costs. Tonight: about \$0.11.

Latency: can it sit inside a real workflow. Tonight: 19 seconds.

